

Continuous-State Dynamic Programming¹

Kenneth L. Judd, Hoover Institution

January 31, 2014

¹Joint work with Yongyang Cai.

Continuous Methods for Continuous-State Problems

- ▶ Basic Bellman equation:

$$V(x) = \max_{u \in D(x)} \pi(u, x) + \beta E\{V(x^+)|x, u\} \equiv (TV)(x).$$

- ▶ Discretization essentially approximates V with a step function
 - ▶ Approximation theory provides better methods to approximate continuous functions.
- ▶ General Task
 - ▶ Find good approximation for V
 - ▶ Identify parameters

General Parametric Approach: Approximating $V(x)$

- ▶ Choose a finite-dimensional parameterization

$$V(x) \doteq \hat{V}(x; a), \quad a \in R^m$$

and a finite number of states

$$X = \{x_1, x_2, \dots, x_n\}$$

- ▶ polynomials with coefficients a and collocation points X
 - ▶ splines with coefficients a with uniform nodes X
 - ▶ rational function with parameters a and nodes X
 - ▶ neural network
 - ▶ specially designed functional forms
- ▶ Objective: find coefficients $a \in R^m$ such that $\hat{V}(x; a)$ “approximately” satisfies the Bellman equation.

General Parametric Approach: Approximating T

For each x_j , $(TV)(x_j)$ is defined by

$$v_j = (TV)(x_j) = \max_{u \in D(x_j)} \pi(u, x_j) + \beta \int \hat{V}(x^+; a) dF(x^+ | x_j, u)$$

and is approximated by \hat{T}

$$v_j = (\hat{T}V)(x_j) \doteq (TV)(x_j)$$

- Integration step: for ω_j and x_j for some numerical quadrature formula

$$\begin{aligned} E\{V(x^+; a) | x_j, u\} &= \int \hat{V}(x^+; a) dF(x^+ | x_j, u) \\ &= \int \hat{V}(g(x_j, u, \varepsilon); a) dF(\varepsilon) \\ &\doteq \sum_{\ell} \omega_{\ell} \hat{V}(g(x_j, u, \varepsilon_{\ell}); a) \end{aligned}$$

- ▶ Maximization step: for $x_i \in X$, evaluate

$$v_i = (T\hat{V})(x_i)$$

- ▶ Hot starts
 - ▶ Concave stopping rules
- ▶ Fitting step:
 - ▶ Data: (v_i, x_i) , $i = 1, \dots, n$
 - ▶ Objective: find an $a \in R^m$ such that $\hat{V}(x; a)$ best fits the data
 - ▶ Methods: determined by $\hat{V}(x; a)$

Approximating T with Hermite Data

Conventional methods just generate data on $V(x_j)$:

$$v_j = \max_{u \in D(x_j)} \pi(u, x_j) + \beta \int \hat{V}(x^+; a) dF(x^+ | x_j, u)$$

- ▶ Envelope theorem:

- ▶ If solution u is interior,

$$v'_j = \pi_x(u, x_j) + \beta \int \hat{V}(x^+; a) dF_x(x^+ | x_j, u)$$

- ▶ If solution u is on boundary

$$v'_j = \mu + \pi_x(u, x_j) + \beta \int \hat{V}(x^+; a) dF_x(x^+ | x_j, u)$$

where μ is a Kuhn-Tucker multiplier

- ▶ Since computing v'_j is cheap, we should include it in data:
 - ▶ Data: (v_i, v'_i, x_i) , $i = 1, \dots, n$
 - ▶ Objective: find an $a \in R^m$ such that $\hat{V}(x; a)$ best fits Hermite data
 - ▶ Methods: determined by $\hat{V}(x; a)$

General Parametric Approach: Value Function Iteration

guess $a \longrightarrow \hat{V}(x; a)$
 $\longrightarrow (v_i, x_i), i = 1, \dots, n$
 $\longrightarrow \text{new } a$

- ▶ Comparison with discretization
 - ▶ This procedure examines only a finite number of states, x_i :
 - ▶ But does *not* assume that the state is always in this finite set.
 - ▶ Choices for the x_i are guided by approximation methods
 - ▶ Procedure examines only a finite number of ε values for the stochastic shocks
 - ▶ But does *not* assume that they are the only ones realized
 - ▶ Choices for the ε_i come from quadrature methods
- ▶ Synergies
 - ▶ Smooth interpolation allows us to use Newton's method for max step.
 - ▶ Smooth interpolation allows more efficient quadrature in (12.7.5).
 - ▶ Efficient quadrature reduces cost of computing objective in max problem

► Finite-horizon problems

- Must use value function iteration since $V(x, t)$ depends on time t .
- Begin with terminal value function, $V(x, T)$
- Compute approximations for each $V(x, t)$, $t = T - 1, T - 2$, etc.

Algorithm 12.5: Parametric Dynamic Programming
with Value Function Iteration

- Objective: Solve the Bellman equation, (12.7.1).
- Step 0: Choose functional form for $\hat{V}(x; a)$, and choose the approximation grid, $X = \{x_1, \dots, x_n\}$.
Make initial guess $\hat{V}(x; a^0)$, and choose stopping criterion $\epsilon > 0$.
- Step 1: Maximization step: Compute
$$v_j = (T\hat{V}(\cdot; a^i))(x_j) \text{ for all } x_j \in X.$$
- Step 2: Fitting step: Using the appropriate approximation method, compute the $a^{i+1} \in R^m$ such that $\hat{V}(x; a^{i+1})$ approximates the (v_i, x_i) data.
- Step 3: If $\|\hat{V}(x; a^i) - \hat{V}(x; a^{i+1})\| < \epsilon$, STOP; else go to step 1.

Convergence

- ▶ T is a contraction mapping
- ▶ \hat{T} may be neither monotonic nor a contraction
- ▶ Shape problems
 - ▶ Standard approximation methods do not preserve shape
 - ▶ Shape problems may become worse with value function iteration

Solution to shape problems

- ▶ Use shape-preserving approximations
 - ▶ Piecewise linear preserves shape in one dimension.
 - ▶ Multilinear approximation does not preserve shape
 - ▶ Shape preserving splines are available for dimensions one and two.
- ▶ Impose shape restrictions in fitting
 - ▶ Use least squares, not interpolation
 - ▶ Add shape constraints to least squares problem
 - ▶ Demand correct slopes at some points
 - ▶ Demand correct curvature at some points.
 - ▶ These methods work well in one dimension, but slow algorithm down considerably for higher dimensions
- ▶ Open research question: What is the best combination of smooth functional form and fitting procedure that preserves shape?

Summary

- ▶ Discretization methods
- ▶ Easy to implement
 - ▶ Numerically stable
 - ▶ Amenable to many accelerations
 - ▶ Poor approximation to continuous problems
- ▶ Continuous approximation methods
 - ▶ Can exploit smoothness in problems
 - ▶ Must work to avoid numerical instabilities
 - ▶ Acceleration is less possible